

exam-2021-questions

February 23, 2021

30E03000 - Data Science for Business I (2021)

1 First exam - 23.02.2021

1.0.1 Case: Healthcare analytics (60pt)

Assume that you are consulting a hospital that aims at improving patient outcomes, efficiency and reducing costs. The data set that the hospital has provided to you contains profiles of several patients including their demographic details along with a set of medical measurements taken during their recent visit (M1 - M9). The objective is to predict which patients should be called for further clinical testing. The testing procedure is costly and the hospital would like to avoid calling healthy patients for an additional test.

The data set is provided in the file **healthcare.csv**. Relevant background information on the variables is provided in the file **healthcare-variables.txt**.

When completing the following steps, you should comment on your findings by adding markdown-boxes along with code-boxes. **Add markdown and code boxes as many as you need.** You can also add subtitles to improve readability of your answer.

Make sure to always answer the questions verbally!! Providing only code will not show your interpretation.

1.0.2 Submission

Download your filled-out notebook in **.ipynb format** and upload it to the MyCourses exam submission box. Name your file studentNum-lastname-firstname-exam.ipynb.

1.1 Task 1 (5pt):

Load data set from “healthcare.csv” file. Perform exploratory data analysis to understand your data better. Check for missing (or otherwise weird) values and variable types. Transform variables when needed (e.g., dummies). You can also drop variables that you don’t want to include, but

justify your decision! Use at least one data visualization technique. Remember to discuss your choices in a markdown box.

```
[17]: import pandas as pd

#add all necessary libraries here
import pandas as pd

#add all necessary libraries here
import numpy as np #scientific computing
import pandas as pd #data management
import itertools

#matplotlib for plotting
import matplotlib.pyplot as plt
from matplotlib import gridspec
import matplotlib.ticker as mtick #for percentage ticks

#sklearn for modeling
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier #Decision Tree algorithm
from sklearn.model_selection import train_test_split #Data split function
from sklearn.preprocessing import LabelEncoder #OneHotEncoding
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc

#Decision tree plot
import pydotplus
from IPython.display import Image

from matplotlib import gridspec
import matplotlib.ticker as mtick #for percentage ticks
import scikitplot as skplt

from sklearn.model_selection import train_test_split #Data split function
from imblearn.over_sampling import SMOTE
from sklearn.decomposition import PCA

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc

from collections import Counter
```

Read the file

```
[32]: # read file into a dataframe
healthcare = pd.read_csv("healthcare.csv")
print(healthcare.head())
```

	Unnamed: 0	sex	age	education	employment	\
0	1	MALE	52	VOCATIONAL/GRAMMAR	LONG-TERM SICK/DISABLED	
1	2	FEMALE	52	SECONDARY	EMPLOYED IN PUBLIC SECTOR	
2	3	FEMALE	78	PRIMARY/NO EDUCATION	OTHER ECONOMICALLY INACTIVE	
3	4	MALE	38	VOCATIONAL/GRAMMAR	EMPLOYED IN PRIVATE SECTOR	
4	5	MALE	61	PRIMARY/NO EDUCATION	RETIRED	

	marital	unempdur	income	abroad	depress	...	m1	m2	m3	m4	\
0	MARRIED	0	560.0	NO	10.0	...	-0.40	2.16	-0.93	1.24	
1	WIDOWED	-8	1600.0	NO	9.0	...	0.92	1.07	0.06	1.19	
2	MARRIED	-8	NaN	NO	10.0	...	1.15	0.71	2.02	-0.18	
3	MARRIED	0	2200.0	NO	1.0	...	1.08	0.34	1.41	-0.99	
4	MARRIED	24	NaN	NO	14.0	...	0.14	0.21	-0.14	1.62	

	m5	m6	m7	m8	m9	y
0	0.74	0.95	0.16	-0.20	3.06	1.0
1	0.03	0.79	0.28	-0.82	0.77	0.0
2	0.92	0.39	0.22	0.85	1.08	0.0
3	0.05	0.07	0.35	1.57	-1.11	0.0
4	-0.33	0.62	0.12	-0.19	1.91	0.0

[5 rows x 25 columns]

```
[58]: print("Dimensions of original data:", data.shape)
```

Dimensions of original data: (1500, 33)

```
[59]: print(list(healthcare.columns))
```

```
['Unnamed: 0', 'sex', 'age', 'education', 'employment', 'marital', 'unempdur',
'income', 'abroad', 'depress', 'trust', 'sport', 'smoke', 'weight', 'bmi', 'm1',
'm2', 'm3', 'm4', 'm5', 'm6', 'm7', 'm8', 'm9', 'y']
```

```
[60]: ## I only keep these variables to do with the model: age, sex, employment,
      ↪ income, bmi, 1-M9 to predict the need for clinical testing
```

```
[61]: modified = healthcare.drop(["Unnamed: 0", "education", "marital", "unempdur",
      ↪ "abroad", "depress", "trust", "sport", "smoke", "weight"], axis=1)
print(modified.columns)
```

```
Index(['sex', 'age', 'employment', 'income', 'bmi', 'm1', 'm2', 'm3', 'm4',
      'm5', 'm6', 'm7', 'm8', 'm9', 'y'],
      dtype='object')
```

```
[62]: modified.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 15 columns):
sex                1500 non-null object
age                1500 non-null int64
employment         1491 non-null object
income             1082 non-null float64
bmi                1483 non-null float64
m1                 1500 non-null float64
m2                 1500 non-null float64
m3                 1500 non-null float64
m4                 1500 non-null float64
m5                 1500 non-null float64
m6                 1500 non-null float64
m7                 1500 non-null float64
m8                 1500 non-null float64
m9                 1500 non-null float64
y                  1498 non-null float64
dtypes: float64(12), int64(1), object(2)
memory usage: 175.9+ KB
```

```
[63]: ## I also drop catagory 'OTHER ECONOMICALLY INACTIVE' and 'UNEMPLOYED' because
      →these two category will not impact the result of clinical testing
```

```
[64]: con1 = modified.employment != "UNEMPLOYED"
      con2 = modified.employment != "OTHER ECONOMICALLY INACTIVE"
      modified[con1 & con2]
```

```
[64]:
```

	sex	age	employment	income	bmi	m1	m2	\
0	MALE	52	LONG-TERM SICK/DISABLED	560.0	22.160665	-0.40	2.16	
1	FEMALE	52	EMPLOYED IN PUBLIC SECTOR	1600.0	26.233556	0.92	1.07	
3	MALE	38	EMPLOYED IN PRIVATE SECTOR	2200.0	26.643599	1.08	0.34	
4	MALE	61	RETIRED	NaN	20.761246	0.14	0.21	
5	MALE	77	RETIRED	1330.0	27.434842	-1.11	0.30	
...	
1493	FEMALE	55	SELF-EMPLOYED	NaN	27.639801	0.63	0.43	
1494	FEMALE	51	LONG-TERM SICK/DISABLED	NaN	33.593750	-0.74	-1.10	
1495	MALE	24	PUPIL OR STUDENT	1600.0	25.432686	0.40	0.73	
1496	FEMALE	52	EMPLOYED IN PRIVATE SECTOR	2000.0	30.811246	1.27	1.16	
1498	FEMALE	53	EMPLOYED IN PUBLIC SECTOR	3000.0	20.820940	-0.66	1.30	

	m3	m4	m5	m6	m7	m8	m9	y
0	-0.93	1.24	0.74	0.95	0.16	-0.20	3.06	1.0
1	0.06	1.19	0.03	0.79	0.28	-0.82	0.77	0.0
3	1.41	-0.99	0.05	0.07	0.35	1.57	-1.11	0.0

```

4    -0.14  1.62 -0.33  0.62  0.12 -0.19  1.91  0.0
5    -0.03 -0.45 -0.08  0.52  0.98  0.31 -0.55  0.0
...
1493  0.77 -0.15  0.79  0.46  0.95 -0.78 -0.54  1.0
1494  0.52 -0.23 -0.54  0.32  0.35  0.89  0.40  0.0
1495 -0.09 -1.24  0.10  0.31  0.73 -0.51  2.03  0.0
1496 -0.47 -1.24  0.46  0.43  0.71 -0.75  0.20  0.0
1498  0.61  0.41  0.17  0.25  0.67 -1.26  1.06  0.0

```

[1237 rows x 15 columns]

```

[65]: modified = modified.dropna(axis=0, how="any")
      modified.info()
      print(modified.head())

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1061 entries, 0 to 1498
Data columns (total 15 columns):
sex                1061 non-null object
age                1061 non-null int64
employment         1061 non-null object
income             1061 non-null float64
bmi                1061 non-null float64
m1                 1061 non-null float64
m2                 1061 non-null float64
m3                 1061 non-null float64
m4                 1061 non-null float64
m5                 1061 non-null float64
m6                 1061 non-null float64
m7                 1061 non-null float64
m8                 1061 non-null float64
m9                 1061 non-null float64
y                  1061 non-null float64
dtypes: float64(12), int64(1), object(2)
memory usage: 132.6+ KB

```

	sex	age	employment	income	bmi	m1	m2	\
0	MALE	52	LONG-TERM SICK/DISABLED	560.0	22.160665	-0.40	2.16	
1	FEMALE	52	EMPLOYED IN PUBLIC SECTOR	1600.0	26.233556	0.92	1.07	
3	MALE	38	EMPLOYED IN PRIVATE SECTOR	2200.0	26.643599	1.08	0.34	
5	MALE	77	RETIRED	1330.0	27.434842	-1.11	0.30	
6	MALE	55	RETIRED	1600.0	27.636054	1.49	2.31	

	m3	m4	m5	m6	m7	m8	m9	y
0	-0.93	1.24	0.74	0.95	0.16	-0.20	3.06	1.0
1	0.06	1.19	0.03	0.79	0.28	-0.82	0.77	0.0
3	1.41	-0.99	0.05	0.07	0.35	1.57	-1.11	0.0
5	-0.03	-0.45	-0.08	0.52	0.98	0.31	-0.55	0.0

```
6 0.63 -0.15 -0.29 0.23 0.65 0.52 -1.03 1.0
```

```
[69]: new = pd.get_dummies(modified[['sex', 'employment']]).head()
modified = pd.concat([new, modified], axis=1)
print(modified.head())
modified.info()
```

```
sex_FEMALE  sex_MALE  employment_EMPLOYED IN PRIVATE SECTOR \
0          0.0      1.0                                0.0
1          1.0      0.0                                0.0
3          0.0      1.0                                1.0
5          0.0      1.0                                0.0
6          0.0      1.0                                0.0
```

```
employment_EMPLOYED IN PUBLIC SECTOR  employment_FARMER \
0                                0.0                0.0
1                                1.0                0.0
3                                0.0                0.0
5                                0.0                0.0
6                                0.0                0.0
```

```
employment_LONG-TERM SICK/DISABLED  employment_OTHER ECONOMICALLY INACTIVE \
0                                1.0                0.0
1                                0.0                0.0
3                                0.0                0.0
5                                0.0                0.0
6                                0.0                0.0
```

```
employment_PUPIL OR STUDENT  employment_RETIRED  employment_SELF-EMPLOYED \
0                                0.0                0.0                0.0
1                                0.0                0.0                0.0
3                                0.0                0.0                0.0
5                                0.0                1.0                0.0
6                                0.0                1.0                0.0
```

```
employment_UNEMPLOYED  sex  age  employment  income \
0          0.0  MALE  52  LONG-TERM SICK/DISABLED  560.0
1          0.0 FEMALE  52  EMPLOYED IN PUBLIC SECTOR 1600.0
3          0.0  MALE  38  EMPLOYED IN PRIVATE SECTOR 2200.0
5          0.0  MALE  77  RETIRED 1330.0
6          0.0  MALE  55  RETIRED 1600.0
```

```
      bmi    m1    m2    m3    m4    m5    m6    m7    m8    m9    y
0  22.160665 -0.40  2.16 -0.93  1.24  0.74  0.95  0.16 -0.20  3.06  1.0
1  26.233556  0.92  1.07  0.06  1.19  0.03  0.79  0.28 -0.82  0.77  0.0
3  26.643599  1.08  0.34  1.41 -0.99  0.05  0.07  0.35  1.57 -1.11  0.0
5  27.434842 -1.11  0.30 -0.03 -0.45 -0.08  0.52  0.98  0.31 -0.55  0.0
6  27.636054  1.49  2.31  0.63 -0.15 -0.29  0.23  0.65  0.52 -1.03  1.0
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1061 entries, 0 to 1498
Data columns (total 26 columns):
sex_FEMALE                    5 non-null float64
sex_MALE                      5 non-null float64
employment_EMPLOYED IN PRIVATE SECTOR  5 non-null float64
employment_EMPLOYED IN PUBLIC SECTOR  5 non-null float64
employment_FARMER             5 non-null float64
employment_LONG-TERM SICK/DISABLED    5 non-null float64
employment_OTHER ECONOMICALLY INACTIVE 5 non-null float64
employment_PUPIL OR STUDENT          5 non-null float64
employment_RETIRED             5 non-null float64
employment_SELF-EMPLOYED          5 non-null float64
employment_UNEMPLOYED           5 non-null float64
sex                             1061 non-null object
age                             1061 non-null int64
employment                     1061 non-null object
income                         1061 non-null float64
bmi                            1061 non-null float64
m1                             1061 non-null float64
m2                             1061 non-null float64
m3                             1061 non-null float64
m4                             1061 non-null float64
m5                             1061 non-null float64
m6                             1061 non-null float64
m7                             1061 non-null float64
m8                             1061 non-null float64
m9                             1061 non-null float64
y                              1061 non-null float64
dtypes: float64(23), int64(1), object(2)
memory usage: 223.8+ KB

```

```

[70]: data = modified.drop(["sex", "employment"], axis =1)
      data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1061 entries, 0 to 1498
Data columns (total 24 columns):
sex_FEMALE                    5 non-null float64
sex_MALE                      5 non-null float64
employment_EMPLOYED IN PRIVATE SECTOR  5 non-null float64
employment_EMPLOYED IN PUBLIC SECTOR  5 non-null float64
employment_FARMER             5 non-null float64
employment_LONG-TERM SICK/DISABLED    5 non-null float64
employment_OTHER ECONOMICALLY INACTIVE 5 non-null float64
employment_PUPIL OR STUDENT          5 non-null float64
employment_RETIRED             5 non-null float64
employment_SELF-EMPLOYED          5 non-null float64

```

```

employment_UNEMPLOYED      5 non-null float64
age                          1061 non-null int64
income                       1061 non-null float64
bmi                           1061 non-null float64
m1                           1061 non-null float64
m2                           1061 non-null float64
m3                           1061 non-null float64
m4                           1061 non-null float64
m5                           1061 non-null float64
m6                           1061 non-null float64
m7                           1061 non-null float64
m8                           1061 non-null float64
m9                           1061 non-null float64
y                            1061 non-null float64
dtypes: float64(23), int64(1)
memory usage: 207.2 KB

```

```
[71]: pd.set_option('display.max_columns', None)
```

```
[72]: data.corr().style
```

```
[72]: <pandas.io.formats.style.Styler at 0x7f7a245b0630>
```

```
[81]: y = data['y'].values #define target variable
      X = data.loc[:, data.columns != 'y'] #define feature matrix
```

1.2 Task 2 (5pt):

Split the data into training (70%) and testing (30%) data sets. Check outcome distributions on training and test datasets.

```
[82]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    ↪random_state = 1234)
```

```
[83]: train_dist = y_train.value_counts() / len(y_train) #normalize absolute count
    ↪values for plotting
      test_dist = y_test.value_counts() / len(y_test)
      data_dist = y.value_counts() / len(y)

      fig, ax = plt.subplots()

      ax.barh(['Test', 'Train', 'Data'], [test_dist[0], train_dist[0], data_dist[0]],
    ↪color='#1f77b4', label='0 (no)')
```



```

ax.barh(['Test','Train','Data'], [test_dist[1], train_dist[1], data_dist[1]],
↳left=[test_dist[0], train_dist[0], data_dist[0]], color='#ff7f0e', label='1
↳(yes)')
ax.set_title('Split visualization')
ax.legend(loc='upper left')
plt.xlabel('Proportion')
plt.ylabel('Partition')

#plot bar values
for part, a, b in zip(['Test', 'Train', 'Data'], [test_dist[0], train_dist[0],
↳data_dist[0]], [test_dist[1], train_dist[1], data_dist[1]]):
    plt.text(a/2, part, str(np.round(a, 2)))
    plt.text(b/2+a, part, str(np.round(b, 2)));

```

```

↳
↳-----
AttributeError                                Traceback (most recent call
↳last)

```

```

<ipython-input-83-e0c838cad8eb> in <module>
----> 1 train_dist = y_train.value_counts() / len(y_train) #normalize
↳absolute count values for plotting
      2 test_dist = y_test.value_counts() / len(y_test)
      3 data_dist = y.value_counts() / len(y)
      4
      5 fig, ax = plt.subplots()

```

AttributeError: 'numpy.ndarray' object has no attribute 'value_counts'

```

[84]: #Define Decision tree classifier with some default parameters
clf = tree.DecisionTreeClassifier(criterion = "gini", random_state = 100,
                                max_depth=3, min_samples_leaf=3)

#Fit the training data
clf.fit(X_train,y_train) #what do we need here?

```

```

↳
↳-----
ValueError                                Traceback (most recent call
↳last)

```

```

<ipython-input-84-f2ee423b2da9> in <module>
    4
    5 #Fit the training data
----> 6 clf.fit(X_train,y_train) #what do we need here?

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
fit(self, X, y, sample_weight, check_input, X_idx_sorted)
    875         sample_weight=sample_weight,
    876         check_input=check_input,
--> 877         X_idx_sorted=X_idx_sorted)
    878     return self
    879

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
fit(self, X, y, sample_weight, check_input, X_idx_sorted)
    147
    148     if check_input:
--> 149         X = check_array(X, dtype=DTYPE, accept_sparse="csc")
    150         y = check_array(y, ensure_2d=False, dtype=None)
    151         if issparse(X):

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
ensure_min_features, warn_on_dtype, estimator)
    576     if force_all_finite:
    577         _assert_all_finite(array,
--> 578                             allow_nan=force_all_finite ==
'allow-nan')
    579
    580     if ensure_min_samples > 0:

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
_assert_all_finite(X, allow_nan, msg_dtype)
    58     msg_err.format
    59     (type_err,
---> 60     msg_dtype if msg_dtype is not None else X.dtype)
    61 )
    62     # for object dtype data, we only check for NaNs (GH-13254)

ValueError: Input contains NaN, infinity or a value too large for
dtype('float32').

```

```
[85]: #Use classifier to predict labels
y_pred = clf.predict(X_test) #what do we need here?
```

```

↳ -----

NotFittedError                                Traceback (most recent call↳
↳ last)

<ipython-input-85-56cee0ea7a3f> in <module>
      1 #Use classifier to predict labels
----> 2 y_pred = clf.predict(X_test) #what do we need here?

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in↳
↳ predict(self, X, check_input)
    416         The predicted classes, or the predict values.
    417         """
--> 418         check_is_fitted(self)
    419         X = self._validate_X_predict(X, check_input)
    420         proba = self.tree_.predict(X)

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in↳
↳ check_is_fitted(estimator, attributes, msg, all_or_any)
    965
    966         if not attrs:
--> 967             raise NotFittedError(msg % {'name': type(estimator).
↳ __name__})
    968
    969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.↳
↳ Call 'fit' with appropriate arguments before using this estimator.
```

```
[86]: #probabilities
y_pred_probs = clf.predict_proba(X_test)
```

```

↳ -----
```

```

NotFittedError                                Traceback (most recent call
↳last)

<ipython-input-86-46af445fb991> in <module>
    1 #probabilities
----> 2 y_pred_probs = clf.predict_proba(X_test)

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳predict_proba(self, X, check_input)
    902             classes corresponds to that in the attribute :term:
↳`classes_`.
    903             """
--> 904         check_is_fitted(self)
    905         X = self._validate_X_predict(X, check_input)
    906         proba = self.tree_.predict(X)

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳check_is_fitted(estimator, attributes, msg, all_or_any)
    965
    966     if not attrs:
--> 967         raise NotFittedError(msg % {'name': type(estimator).
↳__name__})
    968
    969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.
↳Call 'fit' with appropriate arguments before using this estimator.

```

```

[87]: '''
The graphviz library is used to visualize the tree.
'''

#Decision tree plot
import pydotplus
from IPython.display import Image

# Create DOT data
dot_data = tree.export_graphviz(clf, out_file=None,
                                feature_names=X_train.columns,
                                class_names=['need for testing', 'no need for
↳testing'], filled=True) #or use y_train.unique()

# Draw graph

```

```
graph = pydotplus.graph_from_dot_data(dot_data)

# Show graph
Image(graph.create_png())

# Create PNG
#graph.write_png("clf.png") #uncomment this line to save the plot as a .png file
```

```

└─
└─
NotFittedError                                Traceback (most recent call
last)

<ipython-input-87-1d41bfbe3054> in <module>
    10 dot_data = tree.export_graphviz(clf, out_file=None,
    11                                 feature_names=X_train.columns,
--> 12                                 class_names=['need for testing', 'no
need for testing'], filled=True) #or use y_train.unique()
    13
    14 # Draw graph

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_export.py in
export_graphviz(decision_tree, out_file, max_depth, feature_names,
class_names, label, filled, leaves_parallel, impurity, node_ids, proportion,
rotate, rounded, special_characters, precision)
    743     """
    744
--> 745     check_is_fitted(decision_tree)
    746     own_file = False
    747     return_string = False

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
check_is_fitted(estimator, attributes, msg, all_or_any)
    965
    966     if not attrs:
--> 967         raise NotFittedError(msg % {'name': type(estimator).
__name__})
    968
    969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.
Call 'fit' with appropriate arguments before using this estimator.
```

```
#Define Decision tree classifier with some default parameters
clf = tree.DecisionTreeClassifier(criterion = "gini", random_state = 100,
                                max_depth=3, min_samples_leaf=3)

#Fit the training data
clf.fit(X_train,y_train) #what do we need here?
```

```

ValueError                                Traceback (most recent call
↳ last)

<ipython-input-88-f2ee423b2da9> in <module>
    4
    5 #Fit the training data
----> 6 clf.fit(X_train,y_train) #what do we need here?

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳ fit(self, X, y, sample_weight, check_input, X_idx_sorted)
    875         sample_weight=sample_weight,
    876         check_input=check_input,
--> 877         X_idx_sorted=X_idx_sorted)
    878     return self
    879

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳ fit(self, X, y, sample_weight, check_input, X_idx_sorted)
    147
    148     if check_input:
--> 149         X = check_array(X, dtype=DTYPE, accept_sparse="csc")

```

```

150         y = check_array(y, ensure_2d=False, dtype=None)
151         if issparse(X):

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
↳force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
↳ensure_min_features, warn_on_dtype, estimator)
    576         if force_all_finite:
    577             _assert_all_finite(array,
--> 578                             allow_nan=force_all_finite ==
↳'allow-nan')
    579
    580         if ensure_min_samples > 0:

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳_assert_all_finite(X, allow_nan, msg_dtype)
    58             msg_err.format
    59             (type_err,
---> 60             msg_dtype if msg_dtype is not None else X.dtype)
    61         )
    62         # for object dtype data, we only check for NaNs (GH-13254)

ValueError: Input contains NaN, infinity or a value too large for
↳dtype('float32').

```

```

[89]: #Use classifier to predict labels
y_pred = clf.predict(X_test) #what do we need here?

```

```

↳-----

NotFittedError                                Traceback (most recent call
↳last)

<ipython-input-89-56cee0ea7a3f> in <module>
      1 #Use classifier to predict labels
----> 2 y_pred = clf.predict(X_test) #what do we need here?

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳predict(self, X, check_input)
    416         The predicted classes, or the predict values.

```

```

417         """
--> 418         check_is_fitted(self)
419         X = self._validate_X_predict(X, check_input)
420         proba = self.tree_.predict(X)

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳ check_is_fitted(estimator, attributes, msg, all_or_any)
965
966     if not attrs:
--> 967         raise NotFittedError(msg % {'name': type(estimator).
↳ __name__})
968
969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.
↳ Call 'fit' with appropriate arguments before using this estimator.

```

```

[90]: '''
The graphviz library is used to visualize the tree.
'''

#Decision tree plot
import pydotplus
from IPython.display import Image

# Create DOT data
dot_data = tree.export_graphviz(clf, out_file=None,
                                feature_names=X_train.columns,
                                class_names=['NO RESPONSE', 'RESPONSE'],
                                filled=True) #or use y_train.unique()

# Draw graph
graph = pydotplus.graph_from_dot_data(dot_data)

# Show graph
Image(graph.create_png())

# Create PNG
graph.write_png("clf.png") #uncomment this line to save the plot as a .png file

```

```

↳ -----

```



```

NotFittedError                                Traceback (most recent call
↳last)

<ipython-input-90-df6fcba86956> in <module>
    10 dot_data = tree.export_graphviz(clf, out_file=None,
    11                                 feature_names=X_train.columns,
---> 12                                 class_names=['NO RESPONSE',
↳'RESPONSE'], filled=True) #or use y_train.unique()
    13
    14 # Draw graph

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_export.py in
↳export_graphviz(decision_tree, out_file, max_depth, feature_names,
↳class_names, label, filled, leaves_parallel, impurity, node_ids, proportion,
↳rotate, rounded, special_characters, precision)
    743     """
    744
--> 745     check_is_fitted(decision_tree)
    746     own_file = False
    747     return_string = False

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳check_is_fitted(estimator, attributes, msg, all_or_any)
    965
    966     if not attrs:
--> 967         raise NotFittedError(msg % {'name': type(estimator).
↳__name__})
    968
    969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.
↳Call 'fit' with appropriate arguments before using this estimator.

```

```

[91]: importances = clf.feature_importances_
indices = np.argsort(importances)[::-1]
feature_order = np.array([X.columns.values])
i = np.argsort(importances)[::-1]
feature_order = feature_order[:,i]

```

```

↳-----

```

```

NotFittedError                                Traceback (most recent call
↳ last)

<ipython-input-91-aff7d3daae72> in <module>
----> 1 importances = clf.feature_importances_
      2 indices = np.argsort(importances)[::-1]
      3 feature_order = np.array([X.columns.values])
      4 i = np.argsort(importances)[::-1]
      5 feature_order = feature_order[:,i]

/opt/conda/lib/python3.7/site-packages/sklearn/tree/_classes.py in
↳ feature_importances_(self)
    574         (Gini importance).
    575         """
--> 576         check_is_fitted(self)
    577
    578         return self.tree_.compute_feature_importances()

/opt/conda/lib/python3.7/site-packages/sklearn/utils/validation.py in
↳ check_is_fitted(estimator, attributes, msg, all_or_any)
    965
    966     if not attrs:
--> 967         raise NotFittedError(msg % {'name': type(estimator).
↳ __name__})
    968
    969

NotFittedError: This DecisionTreeClassifier instance is not fitted yet.
↳ Call 'fit' with appropriate arguments before using this estimator.

```

[]:

1.3.1 3.1. Analyze the performance of the model. How does it perform on test set? (5pt)

[]:

1.3.2 3.2. Which variables appear to be most important for predicting the success of a campaign? (5pt)

[]:

1.3.3 3.3. Interpret the decision tree produced by the model. What insights can you get from it? (5pt)

[]:

1.4 Task 4 (10pt):

Could the model be improved by using rebalancing techniques? To answer this, check how the model from Task 3 would perform on a balanced dataset. What performance measures are you considering?

[]:

1.5 Task 5 (25pt):

Use the data to train two more competing classification models (in addition to the decision tree models from Task 3 and 4). Justify your choice of settings for the models.

[]:

1.5.1 5.1. Which model has the highest accuracy in testing data? (5pt)

[]:

1.5.2 5.2. Compare the most important predictors of the different models and comment on their similarity and/or difference. Do they make sense? (5pt)

[]:

1.5.3 5.3. Plot ROC graphs to compare model performance. Which of the models would you choose based on training and testing data? Plot additional graphs to evaluate the models and interpret them. (5pt)

[]:

1.5.4 5.4. Suppose that the cost of a single test is 500, and the reward (or costs saved) due to a successfully detected case is 13000. What kind of advice you would give them? You can use various model performance charts to support you recommendation. Note that this is a 10 point question! Verbally explain and justify your recommendation to the hospital stakeholders! (10pt)

[]: